


Figure 1

Schedule matching, within-arm execution repeatability, and cross-arm semantic event alignment answer different questions and require different evidence. Each claim is limited to the declared schedule, trace projection, execution contexts, and event ontology.

Three distinct validation questions




Evidence accumulates left to right; later claims require additional observations rather than stronger wording.

WHAT TO CHECK: Check that no text touches or overlaps the boxes and arrows, and that the three panels read left to right.

Figure 2

Validation and admission flow. Levels 1-7 accumulate scoped evidence. A frozen required failure retains its suppression action; Level 8 cannot repair it. The claim classes are a later conservative reporting taxonomy for this completed case. Level 6 does not itself supply a sampling basis for inferential pairing.

Pairing-assumption protocol and claim admission




Passes are scoped to tested artifacts and contexts; inferential pairing needs a separate sampling or randomization basis.

WHAT TO CHECK: Check that every label fits inside its box, the orange note points cleanly at L7, and the two outcome arrows do not cross anything.

Figure 3

Synthetic failure-detection and admission matrix. P pass, F fail, A admit, D downgrade, dash = blocked gate. The suite demonstrates known logical relations; it does not estimate failure prevalence in external systems.




The event-keyed repair aligns 5/5 shared events; the stateful stream aligns 1/5.

WHAT TO CHECK: Check that the letters are readable in every cell and the legend does not overlap the note below it.

Figure 4

Fixed-battery trace-projection disagreement in the timed-search stress test. Dots show cluster-level disagreement; bars are descriptive empirical reweighting quantiles, not confidence intervals. Neutral labels replace package names.




Dots: cluster-level disagreement share; bars: pooled or stratum-specific empirical reweighting quantiles of the fixed battery (not confidence intervals). All four execution profiles stay together within a cluster.

WHAT TO CHECK: Check that the percentage labels do not collide with the markers and that the note at the bottom is fully visible.

Figure 5

Fixed-battery factorial contrasts with empirical reweighting quantiles for 2,000 schedule-indexed units per cell. Bars are descriptive sensitivities, not confidence intervals or tests.



2,000 schedule-indexed paired units per cell; bars are stratified paired-unit empirical reweighting quantiles of the fixed battery (not confidence intervals). Ranges spanning zero support no effect, equivalence, or superiority conclusion.

WHAT TO CHECK: Check that all four contrast labels and value annotations are readable and that nothing overlaps the dashed zero line.